

trainable params: 4,194,304 || all params: 6,742,609,920 || trainable%: 0.06220594176090199

```
[INFO|trainer.py:641] 2024-07-15 03:48:01,503 >> Using cpu_amp half precision backend
[INFO|trainer.py:2078] 2024-07-15 03:48:01,757 >> ***** Running training *****
[INFO|trainer.py:2079] 2024-07-15 03:48:01,758 >>     Num examples = 990
[INFO|trainer.py:2080] 2024-07-15 03:48:01,760 >>     Num Epochs = 3
[INFO|trainer.py:2081] 2024-07-15 03:48:01,761 >>     Instantaneous batch size per device = 4
[INFO|trainer.py:2084] 2024-07-15 03:48:01,762 >>     Total train batch size (w. parallel, distributed & accumulation) = 8
[INFO|trainer.py:2085] 2024-07-15 03:48:01,763 >>     Gradient Accumulation steps = 2
[INFO|trainer.py:2086] 2024-07-15 03:48:01,765 >>     Total optimization steps = 372
[INFO|trainer.py:2087] 2024-07-15 03:48:01,768 >>     Number of trainable parameters = 4,194,304
```

[372/372 23:19, Epoch 3/3]

Step Training Loss

```
[INFO|trainer.py:2329] 2024-07-15 04:11:52,464 >>
```

Training completed. Do not forget to share your model on huggingface.co/models =)

```
[INFO|trainer.py:3410] 2024-07-15 04:11:52,470 >> Saving model checkpoint to ./tmp
[INFO|tokenization_utils_base.py:2513] 2024-07-15 04:11:52,500 >> tokenizer config file saved in ./tmp/tokenizer_config.json
[INFO|tokenization_utils_base.py:2522] 2024-07-15 04:11:52,505 >> Special tokens file saved in ./tmp/special_tokens_map.json
2024-07-15 04:11:52,520 - finetuning.py - intel_extension_for_transformers.transformers.llm.finetuning.finetuning - INFO - *
** Evaluate After Training**
[INFO|trainer.py:3719] 2024-07-15 04:11:52,526 >> ***** Running Evaluation *****
[INFO|trainer.py:3721] 2024-07-15 04:11:52,526 >>     Num examples = 10
[INFO|trainer.py:3724] 2024-07-15 04:11:52,527 >>     Batch size = 4
```

[3/3 00:00]

***** eval metrics *****

epoch	=	3.0
eval_loss	=	0.0151
eval_ppl	=	1.0152
eval_runtime	=	0:00:01.66
eval_samples	=	10
eval_samples_per_second	=	5.989
eval_steps_per_second	=	1.797